

# Generative Adversarial Networks


Matt Olson

December 1, 2017

# GAN: Overview

- High dimensional density estimation (i.e. unsupervised learning)
- **Generative**: goal is to create *new* samples that resemble training data (see applications)
- **Adversarial**: training proceeds by a “counterfeiter” pitted against discriminator (more on this later)
- **Network**: the “counterfeiter” and discriminator are neural networks
- Yann Lecun “(GANs are) the most interesting idea in the last 10 years in ML”


# Application I: Sampling Images



# Application II: Fashion Generation



# Application III: Image Completion



# Parametric Density Estimation: MLE


Classical MLE:

$$\begin{aligned}\max_{\theta} \prod_{i=1}^n p_{\theta}(x_i) &\Leftrightarrow \min_{\theta} \sum_{i=1}^n \log \frac{1}{p_{\theta}(x_i)} \\ &\Leftrightarrow \min_{\theta} \int \log \left( \frac{\frac{1}{n} \sum_{i=1}^n \delta(x - x_i)}{p_{\theta}(x)} \right) \frac{1}{n} \sum_{i=1}^n \delta(x - x_i) \\ &= KL(p_{data} || p_{\theta})\end{aligned}$$


Maximum likelihood can be viewed as minimizing the KL divergence between a parametric family and the empirical distribution. One might consider other loss functions between these two distributions ...

# Parametric Density Estimation: Other Approaches

Problem setting: generate data from mixture of Gaussians, fit a univariate Gaussian to the data. For sharp samples, may want to consider other loss functions besides negative log-likelihood ...



Maximum likelihood



Reverse KL

# GAN: Shannon-Jensen Divergence

GAN optimization problem:


$$\min_{p_G} KL \left( p_G \parallel \frac{p_G + p_{data}}{2} \right) + KL \left( p_{data} \parallel \frac{p_G + p_{data}}{2} \right)$$
$$\Leftrightarrow$$
$$\min_{p_D} \max_{p_G} -\frac{1}{2} \mathbb{E}_{x \sim p_{data}} \log p_D(x) - \frac{1}{2} \mathbb{E}_{z \sim p_G} \log (1 - p_D(z))$$

where  $p_D$  and  $p_G$  are probability distributions over the input space.

In words,  $p_G$  tries to generate data (images, etc.) that look like the training data, while  $p_D$  tries to distinguish generated data from training data.



# GAN: In Pictures



# GAN: Training Procedure

---

**Algorithm 1** Minibatch stochastic gradient descent training of generative adversarial nets. The number of steps to apply to the discriminator,  $k$ , is a hyperparameter. We used  $k = 1$ , the least expensive option, in our experiments.

---

**for** number of training iterations **do**

**for**  $k$  steps **do**

- Sample minibatch of  $m$  noise samples  $\{z^{(1)}, \dots, z^{(m)}\}$  from noise prior  $p_g(z)$ .
- Sample minibatch of  $m$  examples  $\{x^{(1)}, \dots, x^{(m)}\}$  from data generating distribution  $p_{\text{data}}(x)$ .
- Update the discriminator by ascending its stochastic gradient:

$$\nabla_{\theta_d} \frac{1}{m} \sum_{i=1}^m \left[ \log D(x^{(i)}) + \log (1 - D(G(z^{(i)}))) \right].$$

**end for**

- Sample minibatch of  $m$  noise samples  $\{z^{(1)}, \dots, z^{(m)}\}$  from noise prior  $p_g(z)$ .
- Update the generator by descending its stochastic gradient:

$$\nabla_{\theta_g} \frac{1}{m} \sum_{i=1}^m \log (1 - D(G(z^{(i)}))) .$$


**end for**

The gradient-based updates can use any standard gradient-based learning rule. We used momentum in our experiments.

---

# MNIST: Image Generation


- Sampling digits from MNIST data set
- GAN samples: see `gan.py`
- VAE samples: see `Week5/vae.py`



(a) Sample Digits



(b) VAE Samples



(c) GAN Samples