


Generative Adversarial Networks

Matt Olson

December 1, 2017

GAN: Overview


Application I: Sampling Images



Application II: Fashion Generation



Application III: Image Completion



Parametric Density Estimation: MLE


Classical MLE:

$$\begin{aligned}\max_{\theta} \prod_{i=1}^n p_{\theta}(x_i) &\Leftrightarrow \min_{\theta} \sum_{i=1}^n \log \frac{1}{p_{\theta}(x_i)} \\ &\Leftrightarrow \min_{\theta} \int \log \left(\frac{\frac{1}{n} \sum_{i=1}^n \delta(x - x_i)}{p_{\theta}(x)} \right) \frac{1}{n} \sum_{i=1}^n \delta(x - x_i) \\ &= KL(p_{data} || p_{\theta})\end{aligned}$$


Maximum likelihood can be viewed as minimizing the KL divergence between a parametric family and the empirical distribution. One might consider other loss functions between these two distributions ...

Parametric Density Estimation: Other Approaches

Problem setting: generate data from mixture of Gaussians, fit a univariate Gaussian to the data. For sharp samples, may want to consider other loss functions besides negative log-likelihood ...



Maximum likelihood



Reverse KL

GAN: Shannon-Jensen Divergence


GAN optimization problem:

$$\min_{p_g} KL \left(p_g \parallel \frac{p_g + p_{data}}{2} \right) + KL \left(p_{data} \parallel \frac{p_g + p_{data}}{2} \right)$$
$$\Leftrightarrow$$
$$\min_{p_D} \max_{p_G} -\frac{1}{2} \mathbb{E}_{x \sim p_{data}} \log p_D(x) - \frac{1}{2} \mathbb{E}_{z \sim p_G} \log (1 - p_D(z))$$

where p_D and p_G are probability distributions over the input space.

In words, p_G tries to generate examples that “confuse” the discriminator p_D , and p_D tries to separate items generate from p_D from the data.

GAN: In Pictures



GAN: Training Procedure

Algorithm 1 Minibatch stochastic gradient descent training of generative adversarial nets. The number of steps to apply to the discriminator, k , is a hyperparameter. We used $k = 1$, the least expensive option, in our experiments.

for number of training iterations **do**

for k steps **do**

- Sample minibatch of m noise samples $\{z^{(1)}, \dots, z^{(m)}\}$ from noise prior $p_g(z)$.
- Sample minibatch of m examples $\{x^{(1)}, \dots, x^{(m)}\}$ from data generating distribution $p_{\text{data}}(x)$.
- Update the discriminator by ascending its stochastic gradient:

$$\nabla_{\theta_d} \frac{1}{m} \sum_{i=1}^m \left[\log D(x^{(i)}) + \log (1 - D(G(z^{(i)}))) \right].$$

end for

- Sample minibatch of m noise samples $\{z^{(1)}, \dots, z^{(m)}\}$ from noise prior $p_g(z)$.
- Update the generator by descending its stochastic gradient:


$$\nabla_{\theta_g} \frac{1}{m} \sum_{i=1}^m \log (1 - D(G(z^{(i)}))) .$$

end for

The gradient-based updates can use any standard gradient-based learning rule. We used momentum in our experiments.

MNIST: Image Generation


- Sampling digits from MNIST data set
- GAN samples: see `gan.py`
- VAE samples: see `Week5/vae.py`



(a) Sample Digits



(b) VAE Samples



(c) GAN Samples